

5c. Assignments vs Mutations

Martin Alfaro

PhD in Economics

INTRODUCTION

The upcoming sections will be entirely devoted to **vector mutation**, a process that modifies the contents of a vector. However, to properly cover this subject, we first need to introduce some preliminary concepts, including:


- the distinction between assignment and mutation,
- methods for initializing arrays to eventually mutate them, and
- techniques for vector indexing to select elements.

The current section in particular focuses on **the distinction between assignment and mutation of variables**. The difference between these operations can easily go unnoticed by new users, as both use the operator `=` despite being fundamentally different. Clearly delineating them is important not only for Julia, but also for other programming languages.


SOME BACKGROUND

Recall that **variables** serve as labels for **objects**, where objects in turn hold **values**. Moreover, objects can be classified according to the number of **elements** contained, ranging from scalars (single-element objects such as integers and floating-point numbers) to collections (e.g. vectors).

NUMBER



VECTOR



The distinction between objects and their elements is crucial for the remainder of the section. This is because *assignments apply to objects, whereas mutations apply to elements*. More specifically, assignments rebind variables to new objects, while mutations simply modify existing elements of an existing object.

At first glance, this difference might feel like an implementation detail, without practical implications regarding how you must write code. In Part II of this book, we'll show that performance is heavily influenced by the choice between reassignments and mutations. Updating the elements of an existing vector is typically far more efficient than constructing a brand-new vector with updated values.

DEFINING ASSIGNMENTS AND MUTATIONS


Assignments establish a binding between a variable and an object through the `=` operator. For instance, `x = 3` and `x = [1, 2, 3]` are examples of assignments, in which the objects `3` and `[1, 2, 3]` are respectively bound to `x`.

Mutations, by contrast, **modify the elements of an object, without creating a new one**. These operations also rely on the `=` operator. An example of mutation is given by `x[1] = 0`, which modifies the first element of `x`.


Despite sharing the same operator `=`, assignments and mutations are conceptually distinct. The difference becomes clearer when we **consider objects as residing at specific memory addresses**. Under this view, an assignment for `x` involves two steps: *i*) finding a memory location to store the object, and *ii*) labeling the memory address as `x` for easy access. A mutation, in contrast, modifies the data stored at an existing memory address, leaving the address itself unchanged.

To illustrate, if `x = [6, 7]` is run, `x` becomes associated with an object containing `[6, 7]`. This constitutes an *assignment*. However, if `x[1] = 0` is executed afterward, the operation modifies the original object `[6, 7]`, now becoming `[6, 0]`. This operation constitutes a *mutation*: `x` continues to reference the same memory address, even though its content has changed.

MUTATION



ASSIGNMENT



! Mutating All Elements vs Assignment

Mutating all the elements of `x` doesn't imply a new assignment. For example, modifying `x` using `x[:] = ...` is a mutation.

MUTATING ALL ELEMENTS

```
x = [4,5]
```

```
x[:] = [0,0]
```

```
julia> x
2-element Vector{Int64}:
 0
 0
```

ALIAS VS COPY

Since both assignments and mutations employ the `=` operator, a natural question that arises is when `=` results in an assignment or a mutation.


Next, we focus on **cases where entire objects appear on both sides of `=`**, as in `y = x`. Other scenarios are left for upcoming sections, after we introduce the concept of slices (i.e, subsets of elements from a vector).

When `y = x` is executed, `y` becomes another name for the object referenced by `x`. In other words, `x` and `y` become different labels for the same underlying object. Formally, we say that `y` is an **alias** of `x`.


It's important to stress that `y = x` doesn't bind `y` to `x` itself. Rather, `y` becomes another label for the object that `x` references. This subtle distinction carries a significant practical implication: reassigning `x` to a new object won't affect `y`'s reference.

To clarify this further, let's consider an example where we first execute `x = 2` and then `y = x`. At this point, both `x` and `y` reference the same object, which holds the value `2`. If we eventually execute `x = 4`, the variable `x` will start pointing to a new object holding the value `4`. However, this won't affect the original object that `x` was originally referencing. As a result, `y` will still point to the original object with value `2`.


This behavior is illustrated visually below.

CORRECT**INCORRECT**

CONSEQUENCE



NOT THE CONSEQUENCE



The same conclusion can be observed by examining code execution.

ALIAS VS COPY

```
x = 2    #'x' points to an object with value 2
y = x    #'y' points to the same object as 'x' (do not interpret it as 'y' pointing to 'x')

x = 4    #'x' now points to another object (but 'y' still points to the object holding 2)
```

```
julia> x
4
julia> y
2
```

! **Two variables may contain identical elements and yet refer to different objects**

The claim can be demonstrated using the operators `==` and `===`, which capture two different notions of equality. The expression `x == y` checks whether `x` and `y` have the same *values*, regardless of whether they reference the same object. In contrast, `x === y` checks whether both `x` and `y` point to the same *object*, thus verifying if they point to the same memory address.

By applying these operators, the following example illustrates that objects with identical elements aren't necessarily referencing the same object.

SAME OBJECT

```
x = [4, 5]
```

```
y = x
```

```
julia> x == y
true
```

```
julia> x === y
true
```

SAME ELEMENTS BUT DIFFERENT OBJECT

```
x = [4,5]
```


```
y = [4,5]
```

```
julia> x == y
```


```
true
```

```
julia> x === y
```

```
false
```

GRAPHICAL REPRESENTATION

We've indicated that the operation `y = x` creates a new alias for `x`, turning `y` and `x` two different labels for the same object. This implies that **modifying the elements of either `x` or `y` will necessarily change the elements referenced by both**. The following diagram illustrates this.

GRAPHICAL REPRESENTATION

The corresponding code snippet captures this case.

VARIABLE 'Y' AS AN ALIAS

```
x = [4,5]
```

```
y = x
```

```
x[1] = 0
```

```
julia> x
```

```
2-element Vector{Int64}:
```

```
0
```

```
5
```

```
julia> y
```

```
2-element Vector{Int64}:
```

```
0
```

```
5
```

If you instead aim to treat `x` and `y` as pointing to separate objects, you must use the `copy` function. This creates a *new object* with the same elements as the original. Consequently, any modification to the new object won't affect the original one, allowing you to work with `x` and `y` independently.

VARIABLE 'Y' AS A COPY

```
x = [4, 5]
y = copy(x)
```

```
x[1] = 0
```

```
julia> x
2-element Vector{Int64}:
 0
 5
```

```
julia> y
2-element Vector{Int64}:
 4
 5
```