

9d. Slice Views

Martin Alfaro

PhD in Economics

INTRODUCTION

We previously defined a slice as a subvector derived from a parent vector `x`. Common examples include expressions such as `x[1:2]`, which extracts elements at positions 1 and 2, or `x[x .> 0]`, which selects those elements that are positive. By default, these operations create a copy of the data and therefore allocate memory, except when the slice comprises a single element.

In this section, we address the issue of memory allocations associated with slices. In particular, we highlight the role of **views**, which bypass the need for a copy by directly referencing the parent object. The strategy is effective when slices are indexed through ranges, although it's not for slices that employ Boolean indexing like `x[x .> 0]`, where memory allocation will still occur.

Interestingly, we'll show scenarios where **copying data could actually be faster than using views**, despite the additional memory allocation involved. This apparent paradox emerges because copied data is stored in a contiguous block of memory, which provides more efficient access patterns.

VIEWS OF SLICES

We begin by showing that views don't allocate memory *when a slice is indexed by a range*. This behavior can yield performance improvements over regular slices, which create a copy of the data by default.

SLICE AS A COPY

```
x = [1, 2, 3]
```

```
foo(x) = sum(x[1:2])           # allocations from the slice 'x[1:2]'
```

```
julia> @btime foo($x)
24.838 ns (2 allocations: 80 bytes)
```

SLICE AS A VIEW

```
x = [1, 2, 3]
```

```
foo(x) = sum(@view(x[1:2]))    # it doesn't allocate
```

```
julia> @btime foo($x)
2.924 ns (0 allocations: 0 bytes)
```

Keep in mind, though, that **views created through Boolean indexing neither reduce memory allocations nor are more performant**. This is why you shouldn't rely on views of these objects if your goal is to speed up computations. This fact is illustrated below.

BOOLEAN INDEX (COPY)

```
x = rand(1_000)
```

```
foo(x) = sum(x[x .> 0.5])
```

```
julia> @btime foo($x)
797.000 ns (6 allocations: 4.094 KiB)
```

BOOLEAN INDEX (VIEW)

```
x = rand(1_000)
```

```
foo(x) = @views sum(x[x .> 0.5])
```

```
julia> @btime foo($x)
1.070 μs (6 allocations: 4.094 KiB)
```

COPYING DATA MAY BE FASTER

Although views can reduce memory allocations, there are scenarios where copying data can result in faster performance. A detailed comparison of copies versus views will be provided in another section. Here, we simply remark on this possibility.

Essentially, the choice between copies and views reflects a fundamental trade-off between memory allocation and data access patterns. On one hand, newly created vectors store data in contiguous blocks of memory, enabling more efficient CPU access and allowing for certain optimizations. On the other hand, views avoid allocation, but may also require accessing data scattered across non-contiguous memory regions.

Below, we illustrate a scenario in which the overhead of creating a copy is outweighed by the benefits of contiguous memory access. This makes copying the more efficient choice.

COPY

```
x = rand(100_000)

function foo(x)
    y = x[1:2:length(x)]

    max.(y, 0.5)
end
```

```
julia> @btime foo($x)
61.026 μs (6 allocations: 781.375 KiB)
```

VIEW

```
x = rand(100_000)

function foo(x)
    y = @view x[1:2:length(x)]

    max.(y, 0.5)
end
```

```
julia> @btime foo($x)
364.969 μs (3 allocations: 390.688 KiB)
```